

Deep Learning for Image Analysis

DL4IA – Report for Assignment 1

Student Linus Falk

March 29, 2023

1 Introduction

First assignment in the course Deep learning for image analysis

2 Mathematical exercises

Given the linear regression model:

$$z_i = \sum_{j=1}^p w_j x_{ij} + b \quad (1)$$

with the cost function

$$J = \frac{1}{n} \sum_{i=1}^n L_i \quad (2)$$

where $L_i = (y_i - z_i)^2$

Exercise 1.

$$\frac{\partial J}{\partial w_j} = \frac{1}{n} \sum_{i=1}^n \frac{\partial J}{\partial z_i} \frac{\partial z_i}{\partial w_j}$$

$$\frac{\partial J}{\partial b} = \frac{1}{n} \sum_{i=1}^n \frac{\partial J}{\partial z_i} \frac{\partial z_i}{\partial b}$$

Exercise 2.

$$\begin{aligned} \frac{\partial J}{\partial z_i} &= \frac{1}{n} \sum_{i=1}^n 2(y_i - z_i) \\ \frac{\partial z_i}{\partial b} &= \frac{\partial}{\partial b} \sum_{j=1}^p (w_j x_{ij} + b) = 1 \\ \frac{\partial z_i}{\partial w_j} &= \frac{\partial}{\partial w_j} \sum_{j=1}^p (w_j x_{ij} + b) = \sum_{j=1}^p x_{ij} \end{aligned}$$

3 Code exercises

Exercise 3. *Implement a gradient descent algorithm ...*

Based on the derivations given in exercise 1 and 2, the following functions are implemented to perform gradient descent for linear regression.

- The *initiliaze_parameters* function is used to initialize the weights and offsets by method X.
- The *compute_cost* function takes Y_{pred} and Y as input and calculates the cost according to eq. (2).
- The *model_forward* function takes the input features and compute the predictions with the weight and bias/offset

```
1
2 class NeuralNetwork:
3     # Create a neural network with #hidden layers and #neurons in each layer
4     def __init__(self, features, learningRate, *args):
5         """ param features: number of features in the input
6             param learningRate: set the learningrate/stepsize
7             param *args: number of hidden nodes
8
9             return: Creates a linear regression node for now
10        """
11        self.features = features
12        self.args = args
13        self.learningRate = learningRate
14        self.weights = None
15        self.bias = None
16        self.dW = None
17        self.dB = None
18        self.training_history = []
19
20
21        def initiliaze_parameters(self):
22            """
23                return: returns a weight array with zeros \\
24                according to the input size and a bias variable set to zero
25            """
26
27            self.weights = np.zeros((self.features,1))
28            self.bias = 0
29
30
31            return w, b
32
33        def compute_cost(self, Y_pred, Y):
34            """ param Y_pred: prediction of Y after forward pass
35                param Y: the label
36                ...
37
38                return: the cost
39            """
40            return (Y_pred - Y)
41
42        def model_forward(self, X):
43            """ ...
44                return: the prediction with this set of weight and bias
45            """
46            y_pred = np.dot(X, self.weights) + self.bias
47            return y_pred
48
49        #Forward run wrapper
50        def predict(self, X):
51            return self.model_forward(X)
```

The weights and offsets are optimized to minimize the cost as follows:

- The *train_linear_model* function is used to train a linear model with given input X and labels Y with the given number of iterations.
- The *model_backward* function computes the forward pass first to have a new prediction of Y to calculate the cost. It then computes the gradients.
- The *update_parameters* function takes a step with the calculated gradients and updates the weight and biases.

```

1  #Train the network
2  def train_linear_model(self, X, Y, iterations):
3      """ param X: the input X we want to forward
4          param Y: the label
5          param iteration: number of training iteration
6          ...
7      """
8      for i in range(iterations):
9          nn.model_backward(X, Y)
10         nn.update_parameters()
11
12
13     #Calculate the gradients
14     def model_backward(self, X, Y):
15         """ param X: the input X we want to forward
16             param Y: the label
17             ...
18             return:
19         """
20         samples, features = X.shape
21         Y_pred = self.model_forward(X)
22         cost = self.compute_cost(Y_pred, Y)
23
24         self.dW = (2/samples) * np.dot(X.T, cost)
25         self.dB = (2/samples) * np.sum(cost)
26         self.training_history.append(np.mean(np.abs(cost)))
27
28
29     #Update the weight and bias with the pre-calculated gradients
30     def update_parameters(self):
31         self.weights -= self.learningRate * self.dW
32         self.bias -= self.learningRate * self.dB

```

4 Results

Here follows the mathematical expression (3) of the two trained models. Taking a look at the model with seven inputs we can see that it put more weight into the features: **weight** and **year**. This makes a lot of sense. Since the weight will greatly affect the fuel consumption in a mixed driving situation, accelerating more mass require more fuel. The year is also highly relevant since the car manufacturers has become better and better producing efficient engines.

$$\mathbf{W}_{\text{all features}}^T \mathbf{x} + \mathbf{b}_{\text{all features}} = \begin{bmatrix} -1.20 \\ -0.18 \\ -1.96 \\ -16.84 \\ 0.92 \\ 8.87 \\ 2.65 \end{bmatrix}^T \begin{bmatrix} \text{cylinders} \\ \text{displacement} \\ \text{horsepower} \\ \text{weight} \\ \text{acceleration} \\ \text{year} \\ \text{origin} \end{bmatrix} + [25.63] = [\text{mpg}] \quad (3)$$

$$\mathbf{W}_{\text{one feature}}^T \mathbf{x} + \mathbf{b}_{\text{one feature}} = [-29.03] [\text{horsepower}] + [32.67] = [\text{mpg}]$$

The training history is presented in figure 1. Learning rate 1 was neglected since it is too big of a step and diverges. We can see that the small steps take longer time to converge and that the model with one feature doesn't manage to achieve the same low cost as the model with more features. This is because there is more information in the model with seven features available and can therefore make better predictions. More is not always better though, fitting a model to a feature data with poor resolution that is included can result in worse performance than if that feature was discarded. Taking a look at the predictions of the model with one feature: **horsepower** vs the label **mpg** and plotting the training data we can see that it has fitted the line as expected to the data, see figure 3.

In the case of not normalizing the input we must decrease the step size considerably before it can converge and it will take considerably more steps until we get decent performance in comparison with using normalized data, see figure 2.




Figure 1: Plot showing how the cost (2), evaluated on the training sets, is changing with with number of iterations




Figure 2: Plot showing how the cost (2), evaluated on the **non normalized** training sets, is changing with with number of iterations




Figure 3: Plot showing the model with horsepower as the only input

References

- [1] Asimov, Issac (1942). Runaround